

# Logistic Regression

Andrew Ng

# Classification

Email: Spam / Not Spam?

Online Transactions: Fraudulent (Yes / No)?


Tumor: Malignant / Benign ?

$$y \in \{0, 1\}$$

0: “Negative Class” (e.g., benign tumor)

1: “Positive Class” (e.g., malignant tumor)

$$y \in \{0, 1, 2, 3, \dots, n\}$$



Threshold classifier output  $h_{\theta}(x)$  at 0.5:

If  $h_{\theta}(x) \geq 0.5$ , predict “ $y = 1$ ”

If  $h_{\theta}(x) < 0.5$ , predict “ $y = 0$ ”

Linear regression does not work even with a threshold output

Classification:  $y = 0$  or  $1$

In linear regression  $h_{\theta}(x)$  can be  $> 1$  or  $< 0$

Logistic Regression:  $0 \leq h_{\theta}(x) \leq 1$

# **Logistic Regression**

## **Hypothesis representation**

# Logistic Regression Model


Want  $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = g(\theta^T x)$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Sigmoid function  
Logistic function



# Interpretation of Hypothesis Output

$h_{\theta}(x)$  = estimated probability that  $y = 1$  on input  $x$

Example: If  $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$

$$h_{\theta}(x) = 0.7$$

Tell patient that 70% chance of tumor being malignant

$$h_{\theta}(x) = P(y = 1|x; \theta) \quad \text{“probability that } y = 1, \text{ given } x, \text{ parameterized by } \theta\text{”}$$

$$P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$$
$$P(y = 0|x; \theta) = 1 - P(y = 1|x; \theta)$$

# **Logistic Regression**

## **Decision boundary**



# Logistic Regression

$$h_{\theta}(x) = g(\theta^T x) = p(y = 1|x; \theta)$$


$$g(z) = \frac{1}{1+e^{-z}}$$

Suppose predict “ $y = 1$ ” if  $h_{\theta}(x) \geq 0.5$

$$h_{\theta}(x) = g(\theta^T x) \geq 0.5 \iff \theta^T x \geq 0$$

predict “ $y = 0$ ” if  $h_{\theta}(x) < 0.5$

$$h_{\theta}(x) = g(\theta^T x) < 0.5 \iff \theta^T x < 0$$




$$g(z) \geq 0.5 \iff z \geq 0$$

$$g(z) < 0.5 \iff z < 0$$

# Decision Boundary

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$




$$\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

Predict “ $y = 1$ ” if  $-3 + x_1 + x_2 \geq 0$

$$\theta^T x$$


# Non-linear decision boundaries



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

$$\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

Predict “ $y = 1$ ” if  $-1 + x_1^2 + x_2^2 \geq 0$   
 $x_1^2 + x_2^2 = 1$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$$

# Logistic Regression

## Cost function

Training set:  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

m examples

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \quad x_0 = 1, y \in \{0, 1\}$$



$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameters  $\theta$  ?

# Cost function

Linear regression:  $J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})^2$

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$




# Logistic regression cost function

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Cost = 0 if  $y = 1, h_{\theta}(x) = 1$


But as  $h_{\theta}(x) \rightarrow 0$   
 $\text{Cost} \rightarrow \infty$



Captures intuition that if  $h_{\theta}(x) = 0$ ,  
(predict  $P(y = 1|x; \theta) = 0$ ), but  $y = 1$ ,  
we'll penalize learning algorithm by a very  
large cost.

# Logistic regression cost function

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$





# **Logistic Regression**

## **Simplified cost function and gradient descent**

# Logistic regression cost function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Note:  $y = 0$  or  $1$  always

$$\text{Cost}(h_{\theta}(x), y) = -y \times \log(h_{\theta}(x)) - (1 - y) \times \log(1 - h_{\theta}(x))$$

# Logistic regression cost function

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters  $\theta$ :

$$\min_{\theta} J(\theta)$$

To make a prediction given new  $x$ :

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

# Gradient descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all  $\theta_j$ )

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

# Gradient descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update all  $\theta_j$ )

**Algorithm looks identical  
to linear regression!**

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

vs.

$$h_{\theta}(x) = \theta^T x$$

# **Logistic Regression**

## **Advance optimization**

# Optimization algorithm

Cost function  $J(\theta)$ . Want  $\min_{\theta} J(\theta)$ .

Given  $\theta$ , we have code that can compute

- $J(\theta)$
- $\frac{\partial}{\partial \theta_j} J(\theta)$  (for  $j = 0, 1, \dots, n$ )

Gradient descent:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

# Optimization algorithm

Given  $\theta$ , we have code that can compute

- $J(\theta)$
- $\frac{\partial}{\partial \theta_j} J(\theta)$  (for  $j = 0, 1, \dots, n$ )

Optimization algorithms:

- Gradient descent
- Conjugate gradient
- BFGS
- L-BFGS

Advantages:

- No need to manually pick  $\alpha$
- Often faster than gradient descent.

Disadvantages:

- More complex



# scipy.optimize.fmin\_tnc

**scipy.optimize.fmin\_tnc**(*func*, *x0*, *fprime*=None, *args*=(), *approx\_grad*=0, *bounds*=None, *epsilon*=1e-08, *scale*=None, *offset*=None, *messages*=15, *maxCGit*=-1, *maxfun*=None, *eta*=-1, *stepmx*=0, *accuracy*=0, *fmin*=0, *ftol*=-1, *xtol*=-1, *pgtol*=-1, *rescale*=-1, *disp*=None, *callback*=None)

[\[source\]](#)

Minimize a function with variables subject to bounds, using gradient information in a truncated Newton algorithm. This method wraps a C implementation of the algorithm.

**Parameters:** **func** : callable *func*(*x*, \**args*)

Function to minimize. Must do one of:

1. Return *f* and *g*, where *f* is the value of the function and *g* its gradient (a list of floats).
2. Return the function value but supply gradient function separately as *fprime*.
3. Return the function value and set `approx_grad=True`.

If the function returns None, the minimization is aborted.

**x0** : array\_like

Initial estimate of minimum.

**fprime** : callable *fprime*(*x*, \**args*)

Gradient of *func*. If None, then either *func* must return the function value and the gradient (`f, g = func(x, *args)`) or *approx\_grad* must be True.

**args** : tuple

Arguments to pass to function.

**Returns:****x** : *ndarray*

The solution.

**nfeval** : *int*

The number of function evaluations.

**rc** : *int*

Return code as defined in the RCSTRINGS dict.

**See also:**

[minimize](#) Interface to minimization algorithms for multivariate functions. See the ‘TNC’ *method* in particular.

## Notes

---

The underlying algorithm is truncated Newton, also called Newton Conjugate-Gradient. This method differs from `scipy.optimize.fmin_ncg` in that

1. It wraps a C implementation of the algorithm
2. It allows each variable to be given an upper and lower bound.

**scipy.optimize.minimize**(*fun*, *x0*, *args=()*, *method=None*, *jac=None*, *hess=None*, *hessp=None*, *bounds=None*, *constraints=()*, *tol=None*, *callback=None*, *options=None*) [\[source\]](#)

Minimization of scalar function of one or more variables.

*New in version 0.11.0.*

**Parameters:** **fun** : *callable*

Objective function.

**x0** : *ndarray*

Initial guess.

**args** : *tuple, optional*

Extra arguments passed to the objective function and its derivatives (Jacobian, Hessian).

**method** : *str or callable, optional*

Type of solver. Should be one of

- 'Nelder-Mead'
- 'Powell'
- 'CG'
- 'BFGS'
- 'Newton-CG'
- 'Anneal (deprecated as of scipy version 0.14.0)'
- 'L-BFGS-B'
- 'TNC'
- 'COBYLA'
- 'SLSQP'
- 'dogleg'
- 'trust-ncg'
- custom - a callable object (added in version 0.14.0)

If not given, chosen to be one of `BFGS`, `L-BFGS-B`, `SLSQP`, depending if the problem has constraints or bounds.

Example:

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$$

$$J(\theta) = (\theta_1 - 5)^2 + (\theta_2 - 5)^2$$

$$\frac{\partial}{\partial \theta_1} J(\theta) = 2(\theta_1 - 5)$$

$$\frac{\partial}{\partial \theta_2} J(\theta) = 2(\theta_2 - 5)$$

```
defun cost(theta):  
    return (theta[0]-5)**2 +  
           (theta[1]-5)**2  
  
defun grad(theta):  
    gradient = np.zeros(2);  
    gradient[0] = 2*(theta[0]-5)  
    gradient[1] = 2*(theta[1]-5)  
    return gradient
```

```
initialTheta = np.zeros(2)  
result = opt.fmin_tnc(func=cost, x0=initialTheta,  
                      fprime=grad)  
print(cost(result[0]))
```

$$\text{theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}$$

```
def grad(theta, X, Y)
```

```
    gradient[0] = code to compute  $\frac{\partial}{\partial \theta_0} J(\theta)$ 
```

```
    gradient[1] = code to compute  $\frac{\partial}{\partial \theta_1} J(\theta)$ 
```

```
    :
```

```
    gradient[n] = code to compute  $\frac{\partial}{\partial \theta_n} J(\theta)$ 
```

```
    return gradient
```

```
result = opt.fmin_tnc(func=cost, x0=theta,  
                     fprime=grad, args=(X, Y))
```

# **Logistic Regression**


**Multi-class classification:**

**One-vs-all**


# Multiclass classification

- Email foldering/tagging: Work, Friends, Family, Hobby
- Weather: Sunny, Cloudy, Rain, Snow
- Medical diagrams: Not ill, Cold, Flu

Binary classification:




Multi-class classification:








## One-vs-all (one-vs-rest):



Class 1: 

Class 2: 

Class 3: 



$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta) \quad (i = 1, 2, 3)$$

# One-vs-all

Train a logistic regression classifier  $h_{\theta}^{(i)}(x)$  for each class  $i$  to predict the probability that  $y = i$ .

On a new input  $x$ , to make a prediction, pick the class  $i$  that maximizes

$$\max_i h_{\theta}^{(i)}(x)$$